Working with Large Language Models

University of Hamburg - Spring 2025

Dr. Tobias Vlćek

Welcome to today's workshop!

- Explore Large Language Models (LLMs).
- · You've likely interacted with LLMs already:
 - ChatGPT
 - Google's Gemini
 - Claude
- · Understand:
 - How LLMs work on a high level.
 - What you can do with them.

What is an LLM?

- · LLMs are a type of Artificial Intelligence (AI).
- · Specifically within Natural Language Processing (NLP).
- · Designed to:
 - Understand human language.
 - Generate human language.
 - Interact with human language.
 - In remarkably human-like ways.



Why should you care? LLMs are rapidly transforming how we:

- · Write and edit content
- Develop software
- · Conduct research
- · Analyze data
- · Automate tasks
- · Learn new concepts

Understanding these tools is becoming essential across nearly every professional field.

A Great Overview by 3Blue1Brown

- · Greg Sanderson provides an excellent, concise explanation of LLMs.
- Great starting point to understand LLMs.
- Check out his YouTube channel, 3Blue1Brown for more.

https://www.youtube.com/embed/LPZh9B0jkQs

Today's Key Concepts

- · Fundamental principles behind how LLMs work.
- · The process of training these powerful models.
- · Limitations and ethical considerations of LLMs.
- Practical applications, including pair-programming.
- · How to use LLMs with your own data (Retrieval Augmented Generation).
- · Further resources and tools.

How they work

- · High-level overview of LLM workings.
- · No deep technical details.

Deep Learning and Text Data

- · LLMs are deep learning models.
- · Built using artificial neural networks with many layers ("deep").
- Trained on enormous amounts of text data:
 - Internet
 - Books
 - Articles
 - Code
 - And more!
- Training allows them to learn statistical patterns of language.
- · Become incredibly good at predicting word sequences.
- Foundation for tasks like:
 - **Text Generation:** Creating new text similar to training data.
 - Translation: Converting text between languages.
 - **Summarization:** Condensing text into shorter summaries.
 - Question Answering: Answering questions based on learned information.
 - **Sentiment Analysis:** Determining emotional tone of text.

The Transformer

- Modern LLMs powered by transformer model architecture.
- Introduced in "Attention is All You Need" (Vaswani et al., 2017).
- · Transformers revolutionized NLP.

How Transformers Work (Simplified):

- 1. **Tokenization:** Input text broken into *tokens*.
- 2. **Mathematical Relationships:** Identifies relationships between tokens.
 - Understands word/phrase connections in sentences and context.
- 3. Attention Mechanism: Crucial part.
 - · Focuses on most relevant input parts when generating output.
 - · Pays "attention" to important words for meaning.

Example: The Power of Attention

- Sentence: "The cat sat on the mat because it was warm."
- · Attention mechanism understands:
 - "it" refers to "the mat," not "the cat."

- By considering context of surrounding words and relationships.

Tokens: The Building Blocks

- · LLMs process text as tokens, not whole words.
- · Token can be:
 - Whole word (e.g., "cat")
 - Subword (e.g., "un-", "break", "-able" for "unbreakable")
 - Individual characters

Why use tokens?

- · Handles Unknown Words: Processes unseen words via subword units.
- Efficiency: Learning relationships between tokens is more efficient.

Challenges with Tokenization:

- · Tokenization can be tricky.
- Especially for languages without spaces (Chinese, Japanese).
- Current methods better for languages with similar scripts (English, Spanish, French).
- · Potentially disadvantages other languages.
- · Area of ongoing research.



Think about this: How might this impact the fairness and inclusivity of LLMs?

Context Window: LLM's "Memory"

- · Context window: LLM's short-term memory.
- · Number of tokens the model can process at once.
- · Larger context window:
 - "Remembers" more of conversation/document.
 - More coherent and contextually relevant responses.
 - More computational resources needed.
 - Potentially less coherence if too long.
- · Shorter context window:
 - Faster processing.
 - May lose track of earlier conversation parts.



There is a trade-off between the context window and the size of the model. A larger model can usually handle a larger context window, but it will also take more computational resources to run.

LLM Architectures

- · LLMs are versatile tools.
- · Different tasks need different architectures.
- Main architectures designed for specific strengths.

Encoder-Decoder

- How it works:
 - Encoder: Processes input text -> condensed representation ("thought vector").
 - **Decoder:** Takes representation -> generates output text.
 - Two-step process: understand, then create.
- Analogy: Translation.
 - Understand English sentence (encoder).
 - Construct German sentence (decoder).
- Key Feature: Excellent for tasks with different input and output sequences.

Encoder-Only

- · How it works:
 - Focuses on understanding input.
 - Processes input -> rich representation of meaning.
- Analogy: Detective analyzing crime scene.
 - Gather clues (input) -> build complete picture.
 - No long report needed (output).
- **Key Feature:** Great for classification or information extraction.

Decoder-Only

- · How it works:
 - Focuses on generating text.
 - Takes prompt/conversation history -> predicts next word, then next, etc.
 - Builds output sequence.
- · Analogy: Creative writer or conversation.
 - Writer starts with sentence and continues story.
 - Conversation builds upon previous turns.
- · Key Feature: Ideal for original text creation.

Mixture of Experts (MoE)

· How it works:

- Uses multiple smaller "expert" models.
- Each expert specialized in different aspect.
- "Gating network" decides which expert(s) to use.
- Analogy: Large company with departments.
 - Manager (gating network) assigns project to relevant departments.
- Key Feature: Efficient scaling to large models/datasets.
 - Only subset of experts activated per input.
 Improves performance and reduces cost.

 - Used in large chatbot models.

Training LLMs

• Training an LLM: like teaching a child language, but huge scale.

The Training Process

• LLMs learn via "pre-training" on massive text/code datasets.

1. Input Processing:

- · Model receives token sequence.
- Text broken into tokens (words, subwords, characters).

2. Prediction Task:

- · Given sequence, predict next token.
- Like filling in the blank: "The cat sat on the ___"

3. Learning Loop:

- · Make prediction.
- · Compare to actual answer.
- Adjust internal parameters.
- · Repeat billions of times.

Note

This process requires enormous computational resources - training a large model can cost millions of dollars in computing power!

Computational Power and Parallelism

- Training needs immense computational power.
- · Specialized hardware: GPUs (Graphics Processing Units).
- Model Parallelism: Distributes model parts across multiple GPUs.
 - Speeds up training via parallel processing.
 - Reason for Nvidia's stock increase.

Training Phases

- · Training process often divided into phases:
- · Self-Supervised Learning:
 - Trained on massive dataset without explicit labels.
 - Learns to predict next token.
 - Learns basic language rules and patterns.

- Like learning grammar/vocab by reading books.

Supervised Fine-tuning:

- Trained on smaller, *labeled* dataset for specific tasks.
- Specialized for tasks like question answering, summarization.
- Like specialized course after learning basics.

• Reinforcement Learning from Human Feedback (RLHF):

- Human feedback refines model output.
- Makes output more helpful, honest, and harmless.
- Aligns text generation with human preferences.
- Like teacher feedback to improve writing.

Limitations and Ethical Considerations

- LLMs are powerful but not perfect.
- · Understand limitations and ethical implications.

Bias: A Reflection of the Data

- · LLMs can exhibit biases (gender, racial, cultural, etc.).
- Learn biases from training data.
- · Biased training data -> biased output.

Example:

• Trained mostly on text by men -> reflect male perspectives/stereotypes.

Impact of Bias:

- · Unfair/discriminatory outcomes.
- · Reinforcement of harmful stereotypes.
- · Erosion of trust in Al.

Mitigating Bias:

- **Data Curation:** Carefully select diverse training data. Include data from different demographics, languages, cultures.
- Model Fine-tuning: Train on datasets to counteract bias.
- · Auditing and Evaluation: Rigorously test models for bias.

Adversarial Attacks: Tricking the Model

- · Adversarial attacks: small input changes to mislead LLM.
- · Exploit model vulnerabilities.

Example:

• Slightly change prompt wording -> biased/harmful/nonsensical response.

Why is this a problem?

- · Manipulate LLMs for malicious purposes.
- Spreading misinformation or generating harmful content.

A Broader View

- · LLMs raise broader ethical concerns:
- Job Displacement: Automate human tasks -> potential job losses.

- How to prepare for this shift?
- · Misinformation: Generate realistic fake news easily.
 - How to combat misinformation spread?
- Malicious Use: Deepfakes, harmful content, propaganda.
 - What safeguards are needed?
- Privacy: Trained on personal data.
 - Data usage and privacy regulations needed?
- Accountability: Unclear who is responsible for LLM-generated content.
 - How to hold producers/users accountable?
- Transparency: Difficult to understand LLM decision-making.
 - How to make LLMs more transparent?
- · Addressing ethical concerns requires: researchers, policymakers, public.

Putting LLMs to Work

Explore how to use LLMs practically.

Pair Programming with LLMs

- · Pair programming: two programmers work together.
 - One writes code, other reviews/feedback.
- · LLMs as "Al pair programmer":
 - Code Completion: Suggest code snippets, complete lines.
 - Error Detection: Identify bugs, suggest fixes.
 - Code Generation: Generate functions/blocks from instructions.
 - Code Explanation: Explain how code works.
 - **Best Practices:** Suggest improvements and best practices.

Tools for Pair Programming:

- GitHub Copilot: VS Code integration, Al-powered suggestions.
- Cursor: VS Code fork, more powerful AI features.
- Zed: Code editor for pair programming (LLMs and humans). Local and API models.

Tips for Effective Pair Programming

- Start Small: Begin with code completion, then complex tasks.
- Be Specific: Clear, concise instructions to LLM.
- Review Carefully: Always review LLM-generated code. Mistakes can happen!
- · Learn from the LLM: Learn new techniques and coding styles.



Manage Context: Be mindful of context window. For longer codebases, LLM may "forget" earlier parts. Break down large tasks.

Retrieval Augmented Generation (RAG)

- · Retrieval Augmented Generation (RAG): combines LLMs with external knowledge.
- · LLM with vast library and research assistant.

Why is RAG important?

- Up-to-Date Information: Access current information beyond training data.
- Reduced Hallucinations: Responses grounded in facts, less making things up.
- Domain-Specific Knowledge: Connect to your documents, databases, APIs.
- Expert in your area: Tailored to specific domain without retraining.

How RAG Works:

- 1. **Query:** Ask a question or provide prompt.
- 2. Retrieval: Search external sources for relevant info.
- 3. Augmentation: Add retrieved info to LLM's prompt (context).
- 4. **Generation:** LLM generates response based on:
 - · Pre-trained knowledge.
 - · Retrieved information.

Example:

- · Question: "What are latest advancements in quantum computing?"
- Retrieval: Search scientific publications, news, research databases.
- Augmentation: Key findings added to prompt.
- Generation: Comprehensive, up-to-date answer.

Running LLMs Locally with Ollama

· Running LLMs locally: on your own computer.

Advantages:

- · Privacy: Data stays on your machine.
- · Cost Savings: No API fees.
- Customization: Experiment, fine-tune for needs.
- · Offline Access: Use without internet.
- · Requires understanding hardware and technical setup.
- Ollama: user-friendly tool for local LLMs.

Ollama

- Ollama: free, open-source.
- · Simplifies downloading, installing, running LLMs (macOS, Linux, Windows).
- · Command-line interface (CLI), wide model support.

Key Features of Ollama:

- Easy Installation: Simple download and install.
- Model Management: Download/manage models from library (Hugging Face).
- · Command-Line Interface: Interact via commands.
- API Server: Local API server, integrate in applications.

Hardware Requirements

· Local LLMs need sufficient resources:

Component	Requirement	Notes
RAM	Minimum 8GB	- 16GB recommended- 32GB+ for larger models
CPU	Modern multi-core	- Can run smaller models (< 7B params)- Faster = better performance
GPU	VRAM dependent	- More VRAM = larger models- NVIDIA GPUs preferred
Storage	10GB+ free	- Models can be several GB each

Warning

Performance Note: CPU-only can be much slower compared to GPU.

Billions of Parameters (B)

- · LLMs described by parameters (internal learned values).
- More parameters -> more capable (generally), larger model, more resources.
- Model size not everything, quality depends on data/training/architecture.
- · Newer models can be more efficient.

Size Overview

- Small Models (< 7B): CPU with 16GB+ RAM. Experimentation, less demanding tasks.
- Medium Models (7B 13B): Benefit from GPU, sometimes CPU with 32GB+ RAM.
- Large Models (30B+): Powerful GPU with VRAM needed.
- · Very Large Models (70B+): High-end/multiple GPUs.



Model names like mistral:7b, codellama:34b - "b" = billions of parameters.

Quantization: Making Models Smaller

- · Quantization: Reduce size/compute of LLMs "without significantly impacting performance".
- · Represent parameters with lower precision (e.g., 8-bit, 4-bit integers).

Benefits of Quantization:

- · Smaller Model Size: Less RAM/disk needed.
- Faster Inference: Faster low-precision calculations.
- Lower Power Consumption: Less hardware demand.
- **Trade-offs:** Small accuracy decrease possible. More aggressive quantization = more performance impact.
- Ollama and Quantization: Supports quantized models from Hugging Face.
 - q4_0, q8_0 suffixes indicate quantization level.
 - q4_0: 4-bit (balance size/quality).
 - q8_0: 8-bit (closer to original, larger size).



Quantization levels offer different balance of speed, size, quality. Experiment to find best option for your hardware/needs. Default model often sufficient, quantization not always needed.

Getting Started with Ollama

- 1. Download and Install Ollama: ollama.com.
- 2. Pull a Model: ollama pull mistral:7b (terminal/command prompt).

- Run the Model: ollama run mistral:7b (terminal). Chat interface opens.
 Ask a Question: Type prompt and press Enter.
 Experiment: Try different models from the Ollama website model library. Use Hugging Face for more models/quantization.

Wrap-up

- LLMs are powerful tools for many tasks.Understand their capabilities and limitations.Use them effectively in your work.